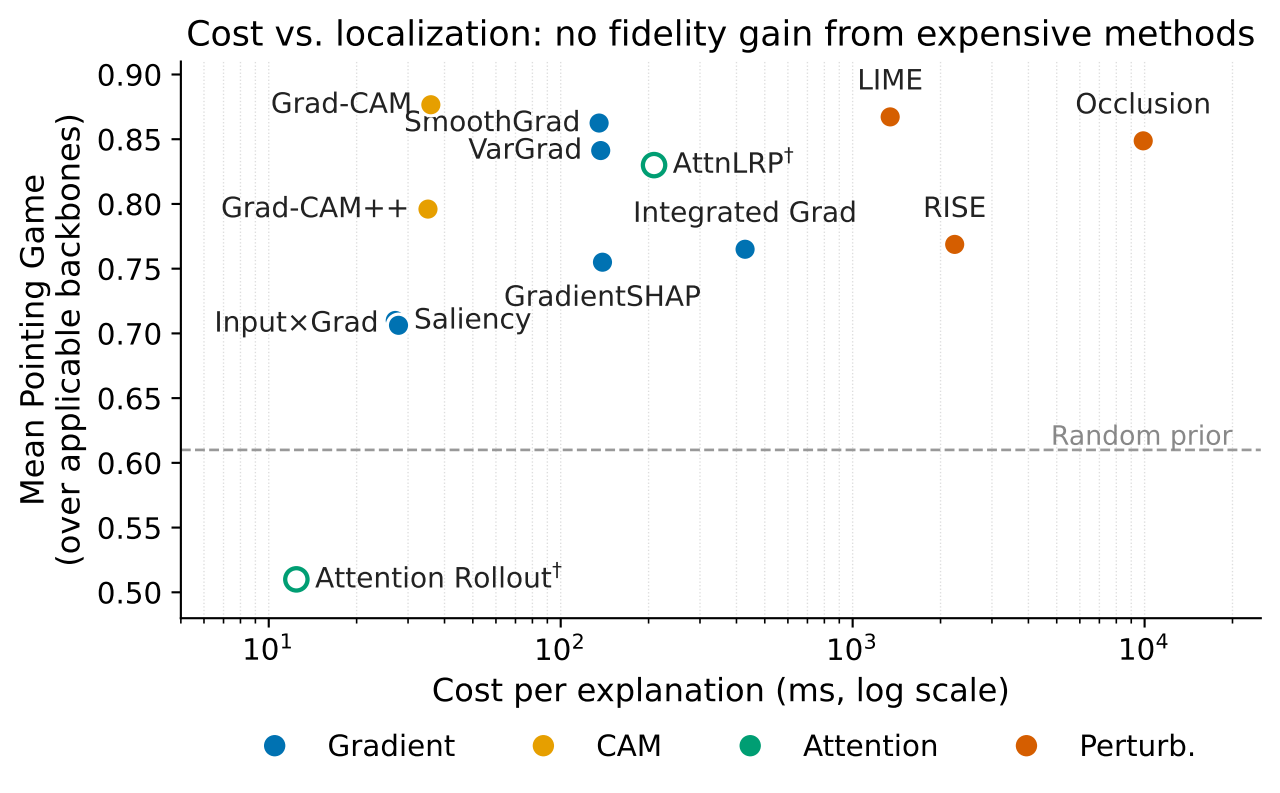
Cost vs. localization: no fidelity gain from expensive methods
Mean Pointing Game
(over applicable backbones)
0.90
0.85
0.80
0.75
0.70
0.65
0.60
0.55
0.50
Grad-CAM
SmoothGrad
VarGrad
AttnLRP†
LIME
Occlusion
Grad-CAM++
Integrated Grad
RISE
GradientSHAP
Input×Grad
Saliency
Random prior
Attention Rollout†
10^1
10^2
10^3
10^4
Cost per explanation (ms, log scale)
Gradient
CAM
Attention
Perturb.